

Cross-attention K/V cache: 1500 encoder positions per layer (the full cache stays in DRAM)

window W_t follows the alignment peak c_t



padding sink S_l : kept exactly, fixed

